



Universität Stuttgart

# Pattern Recognition

## Chapter 11: Probabilistic Discriminative Classifiers

Prof. Dr.-Ing. Uwe Sörgel  
soergel@ifp.uni-stuttgart.de



### Contents

- Generative vs. discriminative classifiers
- Linear Discriminant Function
- Logistic Regression
- Generalized Linear Models
- Training
- Multi-class Problems



Universität Stuttgart



## Generative vs. discriminative classifiers

- **Generative classifiers:**


- Determine the parameters of the likelihood  $p(\mathbf{x}|C)$  (training).
- Determine the prior  $p(C)$ .
- Determine  $p(C|\mathbf{x})$  with the help of the theorem of Bayes.
- **'Generative'**: It is possible to generate synthetic data sets by sampling from the joint distribution  $p(\mathbf{x}, C) = p(\mathbf{x}|C) \cdot p(C)$ .

- **Discriminative classifiers:**


- Direct modelling of  $p(C|\mathbf{x})$
- Focus on the **separating surfaces** in feature space.
- In general, this leads to simpler models and, therefore, requires **fewer training samples**

## Generative and discriminative classifiers

- Example for the comparison of generative and discriminative models for two classes  $C^1$  and  $C^2$ :



This part of  $p(x|C^1)$  has very little influence on the result because  $p(x|C^2)$  is very small.



© Bishop, 2006

The decision based on  $p(C|\mathbf{x})$  considers only the threshold or the separating plane in case of higher dimension.

## Discriminant function

- Often, one is less interested in a model of the probability density than in a sub-division of the feature space into regions which are assigned to the individual classes.
- **Discriminant function:** A function  $g_i(\mathbf{x})$  that assigns  $\mathbf{x}$  to the class  $C^i$  if  $g_i(\mathbf{x}) > g_j(\mathbf{x})$  for all  $i \neq j$ .
- Examples:
  - $p(C^i|\mathbf{x})$
  - $p(\mathbf{x}|C^i) \cdot p(C^i)$
  - $\ln p(\mathbf{x}|C^i) + \ln p(C^i)$
- The discriminant function sub-divides the feature space into regions  $R_i$ , which are assigned to the class  $C^i$ .
- The boundaries of these regions are given by  $g_i(\mathbf{x}) = g_j(\mathbf{x})$ .

## Discriminative methods: Overview

- **Probabilistic discriminative classifiers:**  
Discriminant function is based on  $p(C_i|\mathbf{x})$ 
  - Logistic Regression: binary classification
  - Generalized linear models
- **Non-probabilistic discriminative classifiers:**  
The discriminant function cannot be interpreted as a probability.
  - Decision trees
  - Random forests
  - Boosting
  - Support vector machines
  - Artificial neural networks

## Contents

- Generative vs. discriminative classifiers
- Linear Discriminant Function
- Logistic Regression
- Generalized Linear Models
- Training
- Multi-class Problems

## Discriminant function: linear or non-linear function

- The simplest and most common model is a **linear function** of the input feature vectors  $\mathbf{x}$  of dimension  $D$ :

$$C(\mathbf{w}, \mathbf{x}) = w_0 + w_1 x_1 + \dots + w_D x_D = w_0 + \sum_{i=1}^D w_i x_i = w_0 + \mathbf{w}^T \cdot \mathbf{x}$$

- However, by some modifications we can induce **non-linear** elements too (we will see later why this might be beneficial):

- $M - 1$  non-linear basis functions  $\phi_i$  to **transform input data  $\mathbf{x}$** :

$$C(\mathbf{w}, \mathbf{x}, \Phi) = w_0 + \sum_{i=1}^{M-1} w_i \phi_i(\mathbf{x}) = w_0 + \mathbf{w}^T \cdot \Phi(\mathbf{x})$$

- And/or non-linear transform  $f$  of **outputs**:

$$C(\mathbf{w}, \mathbf{x}, \Phi, f) = f\left(w_0 + \sum_{i=1}^{M-1} w_i \phi_i(\mathbf{x})\right) = f\left(w_0 + \mathbf{w}^T \cdot \Phi(\mathbf{x})\right)$$

## Contents

---

- Generative vs. discriminative classifiers
- Linear Discriminant function
- **Logistic Regression**
- Generalized Linear Models
- Training
- Multi-class Problems

## Logistic sigmoid function

---

- Distinction of two classes  $C^1$ ,  $C^2$  (object, background)

• Theorem of Bayes:

$$p(C^1 | \mathbf{x}) = \frac{p(\mathbf{x} | C^1) \cdot p(C^1)}{p(\mathbf{x} | C^1) \cdot p(C^1) + p(\mathbf{x} | C^2) \cdot p(C^2)}$$
$$= \frac{1}{1 + \frac{p(\mathbf{x} | C^2) \cdot p(C^2)}{p(\mathbf{x} | C^1) \cdot p(C^1)}} = \frac{1}{1 + e^{-a}} := \sigma(a)$$

with  $a(\mathbf{x}) = \ln \frac{p(\mathbf{x} | C^1) \cdot p(C^1)}{p(\mathbf{x} | C^2) \cdot p(C^2)} = \ln \frac{p(C^1 | \mathbf{x})}{p(C^2 | \mathbf{x})}$

Logistic sigmoid function  $\sigma(a) = \frac{1}{1 + e^{-a}}$

## Logistic sigmoid function I

- Originally, this is a generative model, because it is based on the theorem of Bayes.


- $a(\mathbf{x})$  is the negative logarithm of the ratio of the posterior probabilities.

- From now on: **Consideration of  $a(\mathbf{x})$  without Bayesian interpretation.**

- Simple models for  $a(\mathbf{x})$ :  
linear or quadratic functions

- Logistic sigmoid ("S-shaped") function:


$$\sigma(a) = \frac{1}{1 + e^{-a}}$$



## Logistic sigmoid function II

- The **sigmoid function** is useful to model various kinds of processes
  - Growth of population, for example, bacteria:
    - Initially only few individuals, then exponential rise, later slow down and eventually convergence to carrying capacity (the latter defined as the environment's maximal load).
  - Economic process
    - Models diffusion of an innovation through its life cycle (railway, electric power, light bulbs, cars, air travel, transistor, PC, internet, GPS...)
    - Nobel Memorial Prize in Economic Sciences 2000 for Daniel L. McFadden
  - Pattern recognition
    - Neural networks: non-linear *activation function* of an "artificial neuron"

$$\sigma(a) = \frac{1}{1 + e^{-a}}$$



[https://en.wikipedia.org/wiki/Diffusion\\_of\\_innovations](https://en.wikipedia.org/wiki/Diffusion_of_innovations)

## Logistic regression

- (Unrealistic) assumption: The features of  $\mathbf{x}$  are normally distributed with mean values  $\boldsymbol{\mu}_1$  and  $\boldsymbol{\mu}_2$  and identical covariance matrices  $\boldsymbol{\Sigma}_1 = \boldsymbol{\Sigma}_2 = \boldsymbol{\Sigma}$ .

$$\begin{aligned} a(\mathbf{x}) &= \ln \frac{p(\mathbf{x}|C^1) \cdot p(C^1)}{p(\mathbf{x}|C^2) \cdot p(C^2)} = \\ &= -\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_1)^T \cdot \boldsymbol{\Sigma}^{-1} \cdot (\mathbf{x} - \boldsymbol{\mu}_1) + \frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_2)^T \cdot \boldsymbol{\Sigma}^{-1} \cdot (\mathbf{x} - \boldsymbol{\mu}_2) + \ln p(C^1) - \ln p(C^2) = \\ &= \underbrace{(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^T \cdot \boldsymbol{\Sigma}^{-1} \cdot \mathbf{x}}_{\mathbf{w}^T \cdot \mathbf{x}} - \underbrace{\frac{1}{2}\boldsymbol{\mu}_1^T \cdot \boldsymbol{\Sigma}^{-1} \cdot \boldsymbol{\mu}_1 + \frac{1}{2}\boldsymbol{\mu}_2^T \cdot \boldsymbol{\Sigma}^{-1} \cdot \boldsymbol{\mu}_2 + \ln p(C^1) - \ln p(C^2)}_{w_0} = \end{aligned}$$

- Thus  $a(\mathbf{x})$  is a linear function of the features!

$$p(C^1|\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \cdot \mathbf{x} + w_0)}} = \sigma(a(\mathbf{x})) = \sigma(\mathbf{w}^T \cdot \mathbf{x} + w_0)$$

## Logistic regression: Parameters

- In the binary case, we have

$$p(C^2|\mathbf{x}) = 1 - p(C^1|\mathbf{x}) \quad \text{und} \quad 1 - \sigma(a) = \sigma(-a)$$

$$p(C^1|\mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \cdot \mathbf{x} + w_0)}} \quad \text{und} \quad p(C^2|\mathbf{x}) = \frac{1}{1 + e^{(\mathbf{w}^T \cdot \mathbf{x} + w_0)}}$$

- Parameters to be learned:

- Generative view:  $\boldsymbol{\Sigma}, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, p(C^1), p(C^2)$

→ With  $D$  features:  $D \cdot (D + 1) / 2 + 2 \cdot D + 2$  parameters

→ The number of parameters grows quadratically

- Discriminative view:  $\mathbf{w}, w_0$

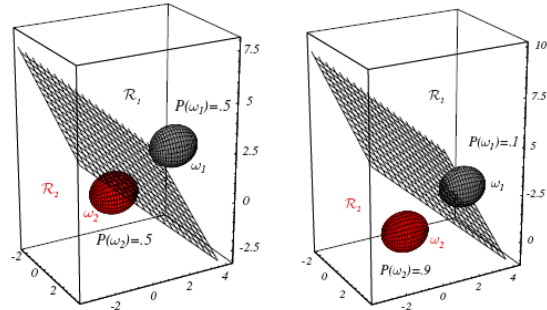
→ With  $D$  features:  $D + 1$  parameters

→ The number of parameters grows linearly

## Logistic regression: Decision boundary

- Class boundary in feature space:  $p(C^1 | \mathbf{x}) = p(C^2 | \mathbf{x})$

$$\begin{aligned} \rightarrow \frac{1}{1 + e^{-(\mathbf{w}^T \cdot \mathbf{x} + w_0)}} &= \frac{1}{1 + e^{(\mathbf{w}^T \cdot \mathbf{x} + w_0)}} \\ \rightarrow -(\mathbf{w}^T \cdot \mathbf{x} + w_0) &= \mathbf{w}^T \cdot \mathbf{x} + w_0 \\ \rightarrow \mathbf{w}^T \cdot \mathbf{x} + w_0 &= 0 \end{aligned}$$



© Duda, Hart, Stork, 2001

- The decision boundary between the classes is a hyperplane.

- Linear discriminative function  $\rightarrow$  Logistic Regression

## Logistic regression: Decision boundary

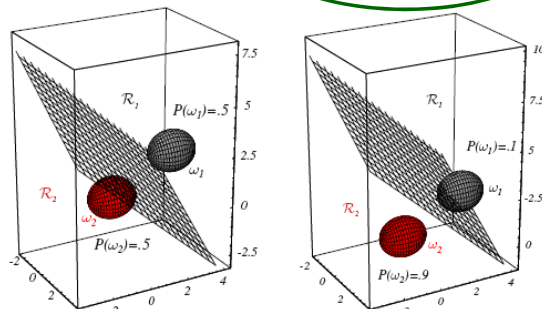
- Class boundary in feature space:  $\mathbf{w}^T \cdot \mathbf{x} + w_0 = 0$

- Normal vector  $\mathbf{w}$ :  $\mathbf{w} = (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2)^T \cdot \boldsymbol{\Sigma}^{-1}$

depends on the vector between the class centers, direction is also influenced by  $\boldsymbol{\Sigma}$ .

- Offset  $w_0$ :  $w_0 = \frac{1}{2} \boldsymbol{\mu}_1^T \cdot \boldsymbol{\Sigma}^{-1} \cdot \boldsymbol{\mu}_1 + \frac{1}{2} \boldsymbol{\mu}_2^T \cdot \boldsymbol{\Sigma}^{-1} \cdot \boldsymbol{\mu}_2 + \ln p(C^1) - \ln p(C^2)$

- Changes to the prior lead to a parallel shift of the decision boundary.



© Duda, Hart, Stork, 2001



## Logistic regression: Geometrical interpretation

• Decision boundary in feature space:  $\varepsilon: \mathbf{w}^T \cdot \mathbf{x} + w_0 = 0$


• For a point  $\mathbf{x}_p$  that does not lie on the separating surface:

$$\mathbf{w}^T \cdot \mathbf{x}_p + w_0 = \|\mathbf{w}\| \cdot d(\varepsilon, \mathbf{x}_p)$$

→ Interpretation of


$$p(C^1 | \mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \cdot \mathbf{x} + w_0)}}$$

as a sigmoid function applied to the (scaled) distance from the separating surface that maps this distance into the interval  $[0, 1]$ !




## Logistic regression: Geometrical interpretation

• Interpretation of  $\|\mathbf{w}\|$ : The larger  $\|\mathbf{w}\|$ , the steeper the sigmoid function.




## Logistic regression: Geometrical interpretation

- Interpretation of  $\|w\|$ : The larger  $\|w\|$ , the steeper the sigmoid function.




## Notion: “Logistic Regression“ I

- „Regression“: The name refers to the initial purpose namely statistical regression. However, it is rather a classification method since we search for an optimal linear separating surface in feature space.
- Different aims of classification, (unsupervised) clustering and regression:



## Notion: “Logistic Regression” II

- „Logistic“: The notion was coined about 1835 from Belgium mathematician Adolphe Quetelet, who desired to discriminate the initial phase of population growth of Belgium from a “logarithmic” (today: exponential) curve:



- The principle that the sigmoid function is applied to a scaled distance to get a probability is often used in other contexts.
- What happens with data that are **not linearly separable**?

## Contents

- Generative vs. discriminative classifiers
- Linear Discriminant function
- Logistic Regression
- **Generalized Linear Models**
- Training
- Multi-class Problems

## Generative Model: Normal distribution with different covariance matrices

- If the covariance matrices are not identical, the quadratic term in the exponent does not disappear:

$$p(C^1 | \mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{x}^T \mathbf{W} \cdot \mathbf{x} + \mathbf{w}^T \cdot \mathbf{x} + w_0)}}$$

with  $\mathbf{W} = 1/2 \cdot (\Sigma_2^{-1} - \Sigma_1^{-1})$


$$\mathbf{w} = \Sigma_1^{-1} \cdot \boldsymbol{\mu}_1 - \Sigma_2^{-1} \cdot \boldsymbol{\mu}_2$$

$$w_0 = 1/2 \cdot \boldsymbol{\mu}_1^T \cdot \Sigma_1^{-1} \cdot \boldsymbol{\mu}_1 + 1/2 \cdot \boldsymbol{\mu}_2^T \cdot \Sigma_2^{-1} \cdot \boldsymbol{\mu}_2 \\ + 1/2 \cdot \ln \|\Sigma_2^{-1}\| - 1/2 \cdot \ln \|\Sigma_1^{-1}\| + \ln p(C^1) - \ln p(C^2)$$

- In general, the class boundary is not a hyperplane but a **hyperquadric**.

## Generative Model: Normal distribution with different covariance matrices

Examples for decision boundaries (3D feature vectors)



© Duda, Hart, Stork, 2001

## Feature space transformations and generalized linear models

• General formula:  $p(C^1 | \mathbf{x}) = \sigma(a(\mathbf{x})) = \frac{1}{1 + e^{-a(\mathbf{x})}}$

- For identical covariance matrices,  $a(\mathbf{x})$  was a linear function of the features  $\mathbf{x}$ :

$$a(\mathbf{x}) = \mathbf{w}^T \cdot \mathbf{x} + w_0$$

- With increasing complexity of the models for the probability densities, the complexity of  $a(\mathbf{x})$  is increased.
  - For example, a quadratic form for normal distributions
- In order still to be able to work with linear functions, one can move on to another feature space:
  - Transformation of the feature space (Feature space mapping)
  - Generalized linear models

## Feature space transformations and generalized linear models

• Feature space mapping  $\Phi(\mathbf{x}) = [\Phi_1(\mathbf{x}), \Phi_2(\mathbf{x}), \dots, \Phi_N(\mathbf{x})]^T$

- $\Phi_i(\mathbf{x})$ : (in principle) arbitrary functions: frequently, polynomials
- $N$ : Dimension of the transformed feature vector (usually greater than the dimension of  $\mathbf{x}$ )
- Frequent choice  $\Phi_1(\mathbf{x}) = 1$

▪ Example for 2D feature space, i.e.  $\mathbf{x} = [x_1, x_2]^T$

$$\Phi(\mathbf{x}) = [1, x_1, x_2, x_1 \cdot x_2, x_1^2, x_2^2]^T$$

- Instead of using a complex model for  $a(\mathbf{x})$ :  
Transition into a higher dimensional feature space in which  $a(\Phi(\mathbf{x}))$  is linear.  
⇒ Generalized linear models

## Feature space transformations and generalized linear models

• Generalized Linear Models: 
$$p(C^1 | \mathbf{x}) = \sigma(a(\mathbf{x})) = \frac{1}{1 + e^{-a(\mathbf{x})}}$$


with  $a(\mathbf{x}) = \mathbf{w}^T \cdot \Phi(\mathbf{x})$

and  $\Phi(\mathbf{x}) = [\Phi_1(\mathbf{x}), \Phi_2(\mathbf{x}), \dots, \Phi_N(\mathbf{x})]^T$ .

- Note: Due to  $\Phi_1(\mathbf{x}) = 1$ ,  $w_0$  becomes the first component of  $\mathbf{w}$ .
- The example of  $\Phi(\mathbf{x}) = (1, x_1, x_2, x_1 \cdot x_2, x_1^2, x_2^2)^T$  leads to a quadratic form for  $a(\mathbf{x})$  similar to the normal distribution!
- Assumptions about the distribution of the features are dropped in favor of a choice of a feature space mapping.



## Examples of feature space mappings I

- In higher dimensional feature space the data can be better separated.



## Examples of feature space mappings II

- In higher dimensional feature space the data can be better separated.



## Examples of feature space mappings III

- Examples of feature space transformations with  $\mathbf{x} = [x_1, x_2]^T$

- Quadratic expansion:  $\Phi(\mathbf{x}) = [1, x_1, x_2, x_1 \cdot x_2, x_1^2, x_2^2]^T$

- Cubic expansion:  $\Phi(\mathbf{x}) = [1, x_1, x_2, x_1 \cdot x_2, x_1^2, x_2^2, x_1^2 \cdot x_2, x_1 \cdot x_2^2, x_1^3, x_2^3]^T$

- Kernel function ("Kernel logistic regression"):

$$\Phi(\mathbf{x}) = [k(\mathbf{x}, \mathbf{x}_1), k(\mathbf{x}, \mathbf{x}_2), \dots, k(\mathbf{x}, \mathbf{x}_N)]^T \quad \text{with kernel function}$$

$$k(\mathbf{x}_n, \mathbf{x}_m) = e^{-\frac{\|\mathbf{x}_n - \mathbf{x}_m\|^2}{2 \cdot \sigma^2}}$$

as a measure for the "distance" of a point from the training points  $\mathbf{x}_n$

- very flexible
- needs hyper-parameter  $\sigma$ .

## Feature space mapping

---

- Using a feature space mapping, linear models can also be applied to problems where the classes are not linearly separable.

- **Disadvantage:** Increase of the number  $N$  of parameters:

- Polynomial expansion: with  $D$  features (incl.  $\Phi_1(\mathbf{x}) = 1$ ), order  $G$ :

$$N = \binom{D+G-1}{G} = \frac{(D+G-1)!}{(D-1)! \cdot G!}$$

- $G = 2 \rightarrow N = D \cdot (D+1) / 2$
    - $G = 3 \rightarrow N = D \cdot (D+1) \cdot (D+2) / 6$
  - Kernel function:  $N$  is equal to the number of training points (we center a kernel at each training point and need to determine a weight for each).
- Could be problematic for feature spaces with  $D > 10$ .

## Contents

---

- Generative vs. discriminative classifiers
- Linear Discriminant function
- Logistic Regression
- Generalized Linear Models
- Training
- Multi-class Problems



## Logistic regression: Training

- **Given:**

- Functional model of feature space mapping
- $N$  points  $\mathbf{x}_i$  with known binary class indicator  $t_i \in \{0, 1\}$
- $t_i$ : **Indicator variable** that shows whether  $\mathbf{x}_i$  belongs to  $C^1$  ( $t_i = 1$ ) or not ( $t_i = 0$ )
- All the indicator variables  $t_i$  can be collected in a vector  $\mathbf{t}$ .

- **Wanted:**

- **Parameter vector**  $\mathbf{w}$  of the generalized linear model

$$p(C^1 | \mathbf{x}) = \frac{1}{1 + e^{-[\mathbf{w}^T \cdot \Phi(\mathbf{x})]}}$$

## Logistic regression: Maximum likelihood training I

- Determine  $\mathbf{w}$  such that  $p(\mathbf{t} | \mathbf{w}, \mathbf{x}_1, \dots, \mathbf{x}_N) \rightarrow \max$  with

$$y_n = p(C^1 | \mathbf{x}_n) = \frac{1}{1 + e^{-[\mathbf{w}^T \cdot \Phi(\mathbf{x}_n)]}} \quad \text{und} \quad p(C^2 | \mathbf{x}_n) = 1 - y_n$$

- Result: 
$$p(\mathbf{t} | \mathbf{w}, \mathbf{x}_1, \dots, \mathbf{x}_N) = \prod_{n=1}^N y_n^{t_n} \cdot (1 - y_n)^{(1-t_n)}$$

because  $p(C^1 | \mathbf{x}_n) > p(C^2 | \mathbf{x}_n) \rightarrow y_n > (1 - y_n)$  for  $t_n = 1$  (i.e.,  $y_n$  contributes) and  
 $p(C^2 | \mathbf{x}_n) > p(C^1 | \mathbf{x}_n) \rightarrow (1 - y_n) > y_n$  for  $t_n = 0$  (i.e.,  $(1 - y_n)$  contributes)

- Instead of the maximization of  $p(\mathbf{t} | \mathbf{w}, \mathbf{x}_1, \dots, \mathbf{x}_N)$ :

- Minimization of the negative log-likelihood:

$$E(\mathbf{w}) = -\ln p(\mathbf{t} | \mathbf{w}, \mathbf{x}_1, \dots, \mathbf{x}_N) \rightarrow \min$$

## Logistic regression: Maximum likelihood training II

- Negative log-likelihood  $E(\mathbf{w})$ :

$$E(\mathbf{w}) = -\sum_{n=1}^N [t_n \cdot \ln(y_n) + (1-t_n) \cdot \ln(1-y_n)] \rightarrow \min$$

- As  $y_n$  depends on  $\mathbf{w}$ ,  $E(\mathbf{w})$  is a non-linear function of  $\mathbf{w}$ .
- Therefore, the minimum of  $E(\mathbf{w})$  can only be determined iteratively.
- Initial values  $\mathbf{w}^0$ : e.g. random numbers
- $E(\mathbf{w})$  is concave and has a single minimum
- Determination of the minimum: **gradient**  $\nabla E(\mathbf{w}) = 0$

## Logistic regression: Maximum likelihood training III

- Newton-Raphson method: using the initial values  $\mathbf{w}^{t-1}$ :

$$\mathbf{w}^t = \mathbf{w}^{t-1} - \mathbf{H}^{-1} \cdot \nabla E(\mathbf{w}^{t-1})$$

- **Gradient**  $\nabla E(\mathbf{w})$ : 
$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \cdot \Phi(\mathbf{x}_n)$$

- Interpretation:  $(y_n - t_n)$  can be considered as classification error for the training point:
  - If  $t_n = 1 \rightarrow C = C^1$  →  $y_n = p(C^1 | \mathbf{x}_n)$  should be close to 1
  - If  $t_n = 0 \rightarrow C = C^2$  →  $y_n$  should be close to 0
- $\nabla E(\mathbf{w})$ : Sum of the feature vectors weighted by  $(y_n - t_n)$

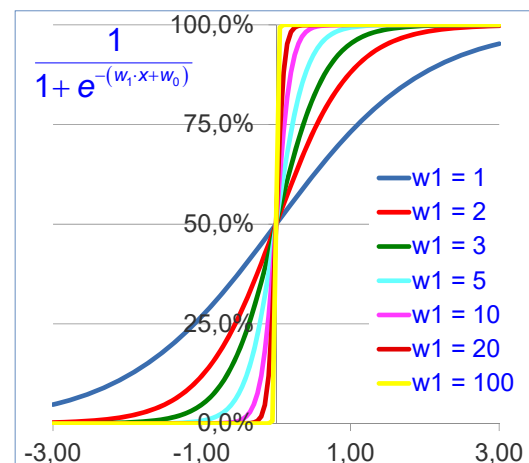
- Hesse-Matrix 
$$\mathbf{H} = \nabla \nabla E(\mathbf{w}) = \sum_{n=1}^N y_n \cdot (1 - y_n) \cdot \Phi(\mathbf{x}_n) \cdot \Phi(\mathbf{x}_n)^T$$

## Maximum likelihood training: Discussion

- Hesse-Matrix is positive definite  $\rightarrow$  inverse exists
- In order to avoid numerical problems:
  - **Scaling of the features:**
    - Determination of mean value  $\mu$  and standard deviation  $\sigma$  of all features from the training data.
    - Shift by  $\mu$ , scaling with  $1 / \sigma$   
 $\rightarrow$  Features all have the same range of values
    - The same scaling has to be applied for training and classification!
  - **ML has the tendency to overfit** the classifier to the training data:
    - For example,  $|\mathbf{w}|$  might become very large  $\rightarrow$  sigmoid approximates step function!
    - ML provides no means to enforce certain desired behavior (i.e., apply knowledge!)  
 $\rightarrow$  **Bayesian method**, which here is **equivalent to regularization** with prior for  $\mathbf{w}$ .

## Logistic regression: Training with regularization I

- MAP: Maximization of  $p(\mathbf{w} | \mathbf{t}, \mathbf{x}_1, \dots, \mathbf{x}_N) = p(\mathbf{t} | \mathbf{w}, \mathbf{x}_1, \dots, \mathbf{x}_N) \cdot p(\mathbf{w})$
- $p(\mathbf{t} | \mathbf{w}, \mathbf{x}_1, \dots, \mathbf{x}_N)$  corresponds to the Likelihood (as with ML)
- **Prior  $p(\mathbf{w})$ :**
  - Sigmoid slope depends on the size of the numerical values of the coefficients  $w_i$  in  $\mathbf{w}$ :
    - The larger  $|w_i|$ , the steeper the sigmoid function.
    - For  $w_i \rightarrow \infty$  the sigmoid function becomes a step function.



## Logistic regression: Training with regularization II

$$p(\mathbf{w} | \mathbf{t}, \mathbf{x}_1, \dots, \mathbf{x}_N) = p(\mathbf{t} | \mathbf{w}, \mathbf{x}_1, \dots, \mathbf{x}_N) \cdot p(\mathbf{w}) \rightarrow \max$$

- To keep the numerical values of  $\mathbf{w}$  small:
- Prior  $p(\mathbf{w})$ : Normal distribution with expectation value  $\mathbf{0}$  and covariance matrix  $\sigma^2 \cdot \mathbf{I}$
- Corresponds to regularization in adjustment theory.
- Requires hyper-parameter  $\sigma$ , which is either fixed by the user or determined via a procedure such as cross-validation.
- Negative logarithm (excluding constant terms):

$$E(\mathbf{w}) = -\sum_{n=1}^N [t_n \cdot \ln(y_n) + (1-t_n) \cdot \ln(1-y_n)] + \frac{\mathbf{w}^T \cdot \mathbf{w}}{2 \cdot \sigma^2} \rightarrow \min$$

## Logistic regression: Training with regularization III

- Minimization of

$$E(\mathbf{w}) = -\sum_{n=1}^N [t_n \cdot \ln(y_n) + (1-t_n) \cdot \ln(1-y_n)] + \frac{\mathbf{w}^T \cdot \mathbf{w}}{2 \cdot \sigma^2} \rightarrow \min$$

leads to the numerical values of  $\mathbf{w}$  that are as small as possible.

- The gradient has to be extended compared to the ML method :

$$\nabla E(\mathbf{w}) = \sum_{n=1}^N (y_n - t_n) \cdot \Phi(\mathbf{x}_n) + \frac{1}{\sigma^2} \cdot \mathbf{w}$$


- This is also true for the Hesse Matrix:

$$\mathbf{H} = \nabla \nabla E(\mathbf{w}) = \sum_{n=1}^N [y_n \cdot (1-y_n) \cdot \Phi(\mathbf{x}_n) \cdot \Phi(\mathbf{x}_n)^T] + \frac{1}{\sigma^2} \cdot \mathbf{I}$$


This means, in the main diagonal, the weights of the direct observations for  $\mathbf{w}$  are added (as in the case of regularization in adjustment).

## Logistic Regression: Example (ML-Training)

Two classes, two features: linearly separable case



Training samples in  
feature space (800)



Regions assigned to the two classes  
in feature space

## Logistic regression: Example

Two classes, two features: linearly separable case  
posterior probabilities



$$p(C^1|x_1,x_2)$$

white ... high probability, black ... low probability



$$p(C^2|x_1,x_2)$$


The sharp boundaries suggest a higher security of the classification in the border region than is actually achieved → Overfitting!

## Logistic regression: Example




Two classes, two features: non linearly separable case

→ The classifier



Training samples in  
feature space (800)



Regions assigned to the two classes  
in feature space


## Logistic regression: Example

Two classes, two features: non-linearly separable case




$p(C^1|x_1,x_2)$

white ... high probability,  
black ... low probability



$p(C^2|x_1,x_2)$

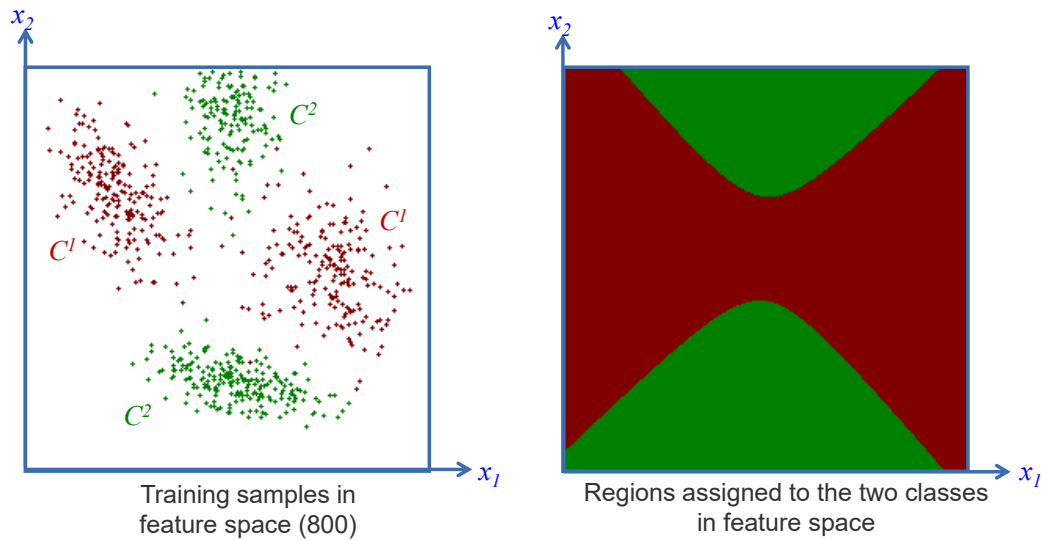
- Small differences in the posterior probabilities
- Relatively large value for  $E(\mathbf{w})$



log-likelihood as a function of  
the iteration count in training

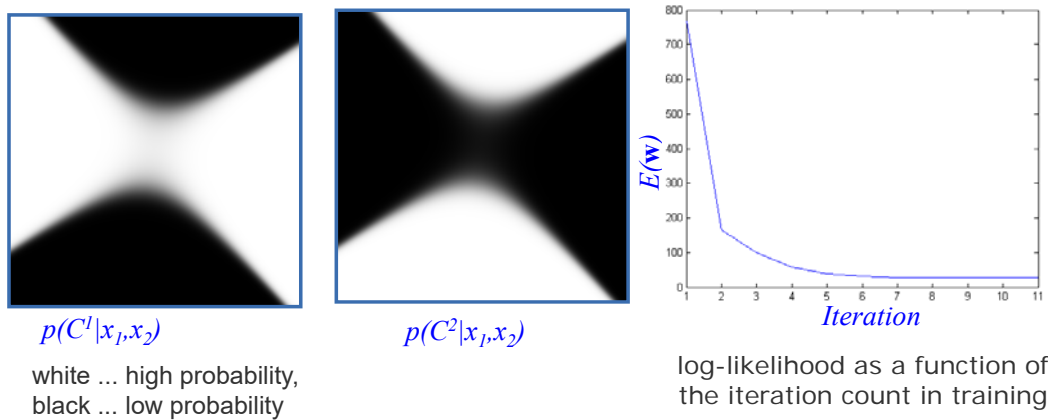
## Logistic regression: Example

Two classes, two features: non-linearly separable case  
 Feature space transformation: the classes can be separated!



## Logistic regression: Example

Two classes, two features: non-linear separated case  
 with feature space transform



- Significant differences in the posterior probabilities
- Low value for  $E(\mathbf{w})$  is reached

## Contents

---

- Generative vs. discriminative classifiers
- Linear Discriminant function
- Logistic Regression
- Generalized Linear Models
- Training
- Multi-class Problems

## Transition to multi-class problems

---



- The posterior probability  $p(C^k|\mathbf{x})$  for each class  $C^k$  can be modelled using the

$$p(C^k|\mathbf{x}) = \frac{e^{a_k(\mathbf{x})}}{\sum_j e^{a_j(\mathbf{x})}}$$

$$\text{with } a_k(\mathbf{x}) = \ln p(\mathbf{x}|C^k) + \ln p(C^k)$$

- Assumptions about  $p(\mathbf{x}|C^k)$  and  $p(C^k)$  lead to models for  $a_k(\mathbf{x})$
- Again, feature space mapping can help to obtain linear models:

$$a_k(\mathbf{x}) = a_k(\Phi(\mathbf{x})) = \mathbf{w}_k^T \cdot \Phi(\mathbf{x})$$

- In training, one parameter vector  $\mathbf{w}_k$  per class has to be determined.



## Multi-class logistic regression: Training

• Softmax function: 
$$p(C^k | \mathbf{x}_n) = \frac{\exp[\mathbf{w}_k^T \cdot \Phi(\mathbf{x}_n)]}{\sum_{j=1}^M \exp[\mathbf{w}_j^T \cdot \Phi(\mathbf{x}_n)]} = y_{nk}$$

• **Training:** Class label  $C_n$  is given for each training point  $\mathbf{x}_n$

• Maximum Likelihood training is similar to the two-class case: the negative log-likelihood has to be minimized:

$$E(\mathbf{w}_1, \dots, \mathbf{w}_M) = - \sum_{n=1}^N \sum_{k=1}^M t_{nk} \cdot \ln(y_{nk}) \rightarrow \min$$

with the binary indicator variables  $t_{nk} = \begin{cases} 1 & \text{if } C_n = C^k \\ 0 & \text{otherwise} \end{cases}$

$M$ ... Number of classes

## Multi-class logistic regression: Maximum likelihood training

• Again, the the Newton-Raphson can be applies: Using the current values from the previous iteration, the weights are updated according to  $\mathbf{w}^{\tau-1}$  from the previous iteration, the weights are updated according to

$$\mathbf{w}^{\tau} = \mathbf{w}^{\tau-1} - \mathbf{H}^{-1} \cdot \nabla E(\mathbf{w}^{\tau-1})$$

• The parameter vectors are not independent

→ **One parameter vector must be declared to be constant,**

e.g.  $\mathbf{w}_1^T = (0, \dots, 0)^T$

•  $\mathbf{w}_1$  is not changed in the optimization procedure

→ The parameter vector  $\mathbf{w}$  to be determined if  $M$  classes are to be discerned becomes:  $\mathbf{w} = (\mathbf{w}_2^T, \dots, \mathbf{w}_M^T)^T$

## Multi-class logistic regression: Maximum likelihood training

- Gradient of the negative log-likelihood

(Derivative of  $E$  by the weight vector of the class  $j$ ):

$$\nabla_{\mathbf{w}_j} E(\mathbf{w}_1, \dots, \mathbf{w}_M) = \sum_{n=1}^N (y_{nj} - t_{nj}) \cdot \Phi(\mathbf{x}_n)$$

- Total gradient vector:

$$\nabla E(\mathbf{w}_1, \dots, \mathbf{w}_M) = \left[ \nabla_{\mathbf{w}_1} E(\mathbf{w}_1, \dots, \mathbf{w}_M)^T, \dots, \nabla_{\mathbf{w}_M} E(\mathbf{w}_1, \dots, \mathbf{w}_M)^T \right]^T$$

- Again, the gradient can be interpreted as the sum of the (transformed) feature vectors weighted by the “classification error”  $(y_{nj} - t_{nj})$ .

## Multi-class logistic regression: Maximum likelihood training

- Hesse matrix  $\mathbf{H}$  also consists of several components:

$$\mathbf{H} = \begin{pmatrix} \mathbf{H}_{22} & \mathbf{H}_{23} & \cdots & \mathbf{H}_{2M} \\ \mathbf{H}_{23}^T & \mathbf{H}_{33} & \cdots & \mathbf{H}_{3M} \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{H}_{2M}^T & \mathbf{H}_{3M}^T & \cdots & \mathbf{H}_{MM} \end{pmatrix}$$

with

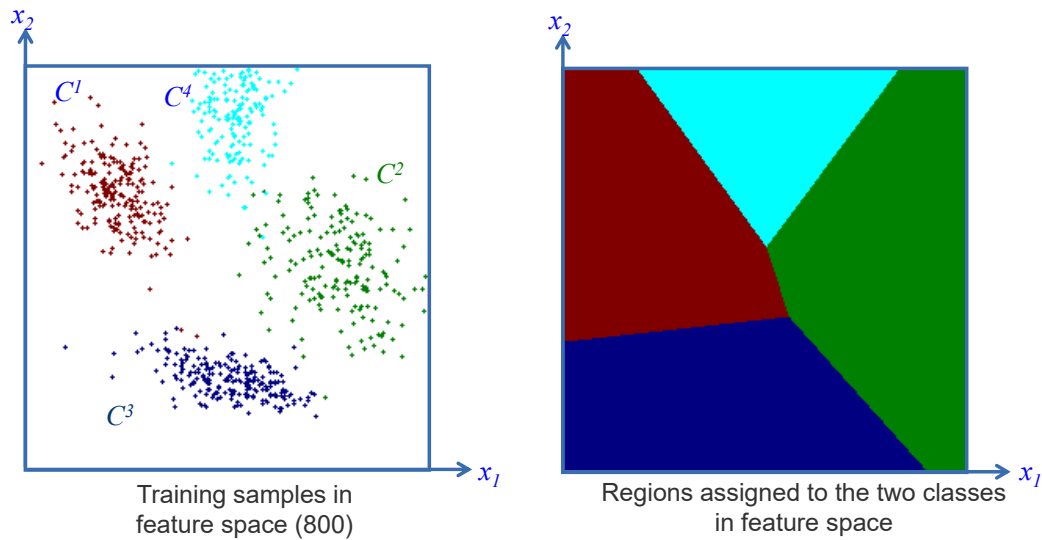
$$\mathbf{H}_{jk} = \nabla_{\mathbf{w}_j} \nabla_{\mathbf{w}_k} E(\mathbf{w}) = \sum_{n=1}^N y_{nk} \cdot (\mathbf{I}_{nk} - y_{nj}) \cdot \Phi(\mathbf{x}_n) \cdot \Phi(\mathbf{x}_n)^T$$

$\mathbf{I}_{nk}$  ... Elements of a unit matrix

- Regularisation: As in the binary case (Gaussian prior with expectation  $\mathbf{0}$  and Covariance  $\sigma \cdot \mathbf{I}$ )

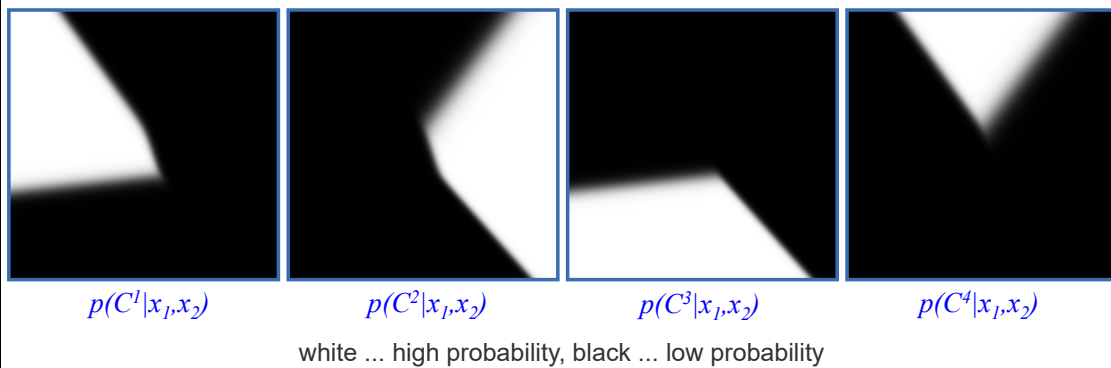
## Multi-class logistic regression: Example (ML-Training)

Four classes, two features



## Multi-class logistic regression: Example (ML-Training)


Four classes, two features:  
posterior probabilities



In the areas where the feature distributions overlap,  
the boundaries are slightly blurred

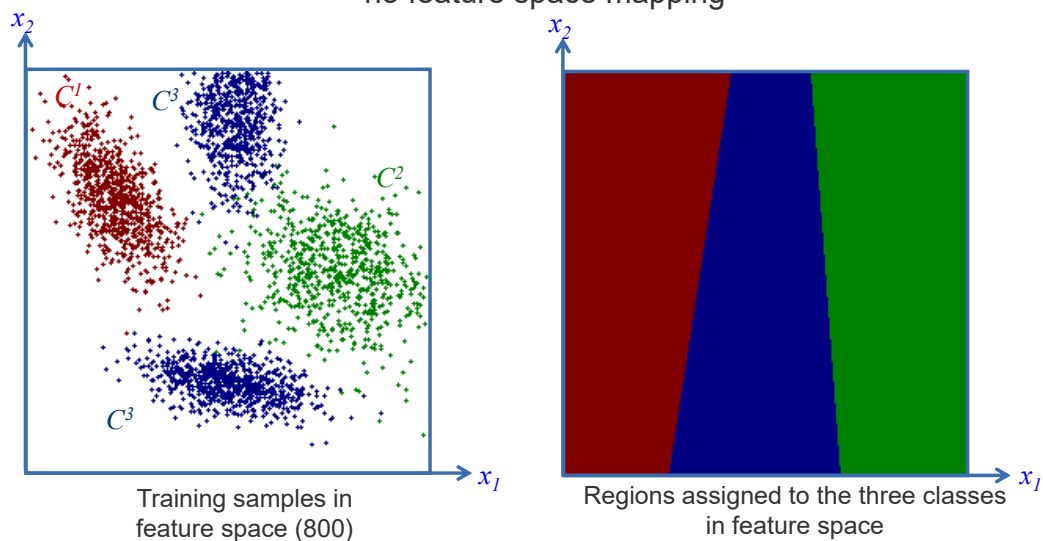
## Multi-class logistic regression: Example (ML-Training)

Four classes, two features:  
Development of the log-likelihood during training




## Multi-class case: Example (ML-Training)

Three classes, two features, not linearly separable  
no feature space mapping




## Multi-class case: Example (ML-Training)

Three classes, two features, not linearly separable, no feature space mapping:  
development of log-likelihood during training



## Multi-class case: Example (ML-Training)

Three classes, two features, not linearly separable –  
quadratic expansion



## Multi-class case: Example (ML-Training)


Three classes, two features, not linearly separable –  
quadratic expansion: posterior probabilities



$p(C^1|x_1,x_2)$



$p(C^2|x_1,x_2)$




$p(C^3|x_1,x_2)$

white ... high probability, black ... low probability

- In the areas where the feature distributions overlap, the boundaries are slightly blurred.
- However, in general there is a very clear distinction → **Overfitting**


## Multi-class case: Example (ML-Training)

Three classes, two features, not linearly separable –  
quadratic expansion: development of log-likelihood during training




## Multi-class case: Example (ML-Training)

Three classes, two features, not linearly separable – quadratic expansion, training with relatively strong regularization ( $\sigma = 2$ )



## Multi-class case: Example (ML-Training)

Three classes, two features, not linearly separable – quadratic expansion, training with regularization: posterior probabilities



white ... high probability, black ... low probability

- Much smoother transitions, uncertainty of the classification is better represented
- Class boundaries may be regularized too strongly

## Discussion

---

- **Discriminative probabilistic methods** directly model the posterior probability
  - No assumption about the distribution of data required
  - Basically, **boundaries between classes are learned**
  - **Linear Models** with / without feature space transformation
    - Fewer parameters to be determined
    - Fewer training data is required
  - Can be expanded to multi-class problems
  - Efficient learning / classification
  - Probabilistic output simplifies further processing.

## Discussion

---

- Despite feature space transformation, the functional model can not fit properly to the distribution of the data
  - **Transition to non-probabilistic methods**
- High-dimensional feature vectors can lead to a large number of parameters to be learned.
- Numerical problems → scaling of the features in training and during the classification
- **ML-Learning**: Problem of overfitting → Regularisation
  - Requires prior for the parameter vector  $\mathbf{w}$ 
    - Hyper-parameter  $\sigma$  (cross validation)



## Literatur

---

- Bishop, C. : Pattern Recognition and Machine Learning. 1st edition, Springer, New York, USA, 2006.
- Duda, R. O., Hart, P. E., Stork, D. G.: Pattern Classification. 2nd edition, Wiley & Sons, New York, USA, 2001.
- Rottensteiner, Franz, 2014: Skript Bildanalyse II, IPI, Leibniz Universität Hannover